


# Distributed Systems and the Web

**Karl Kirch**  
@joekarl



**SIMPLE**

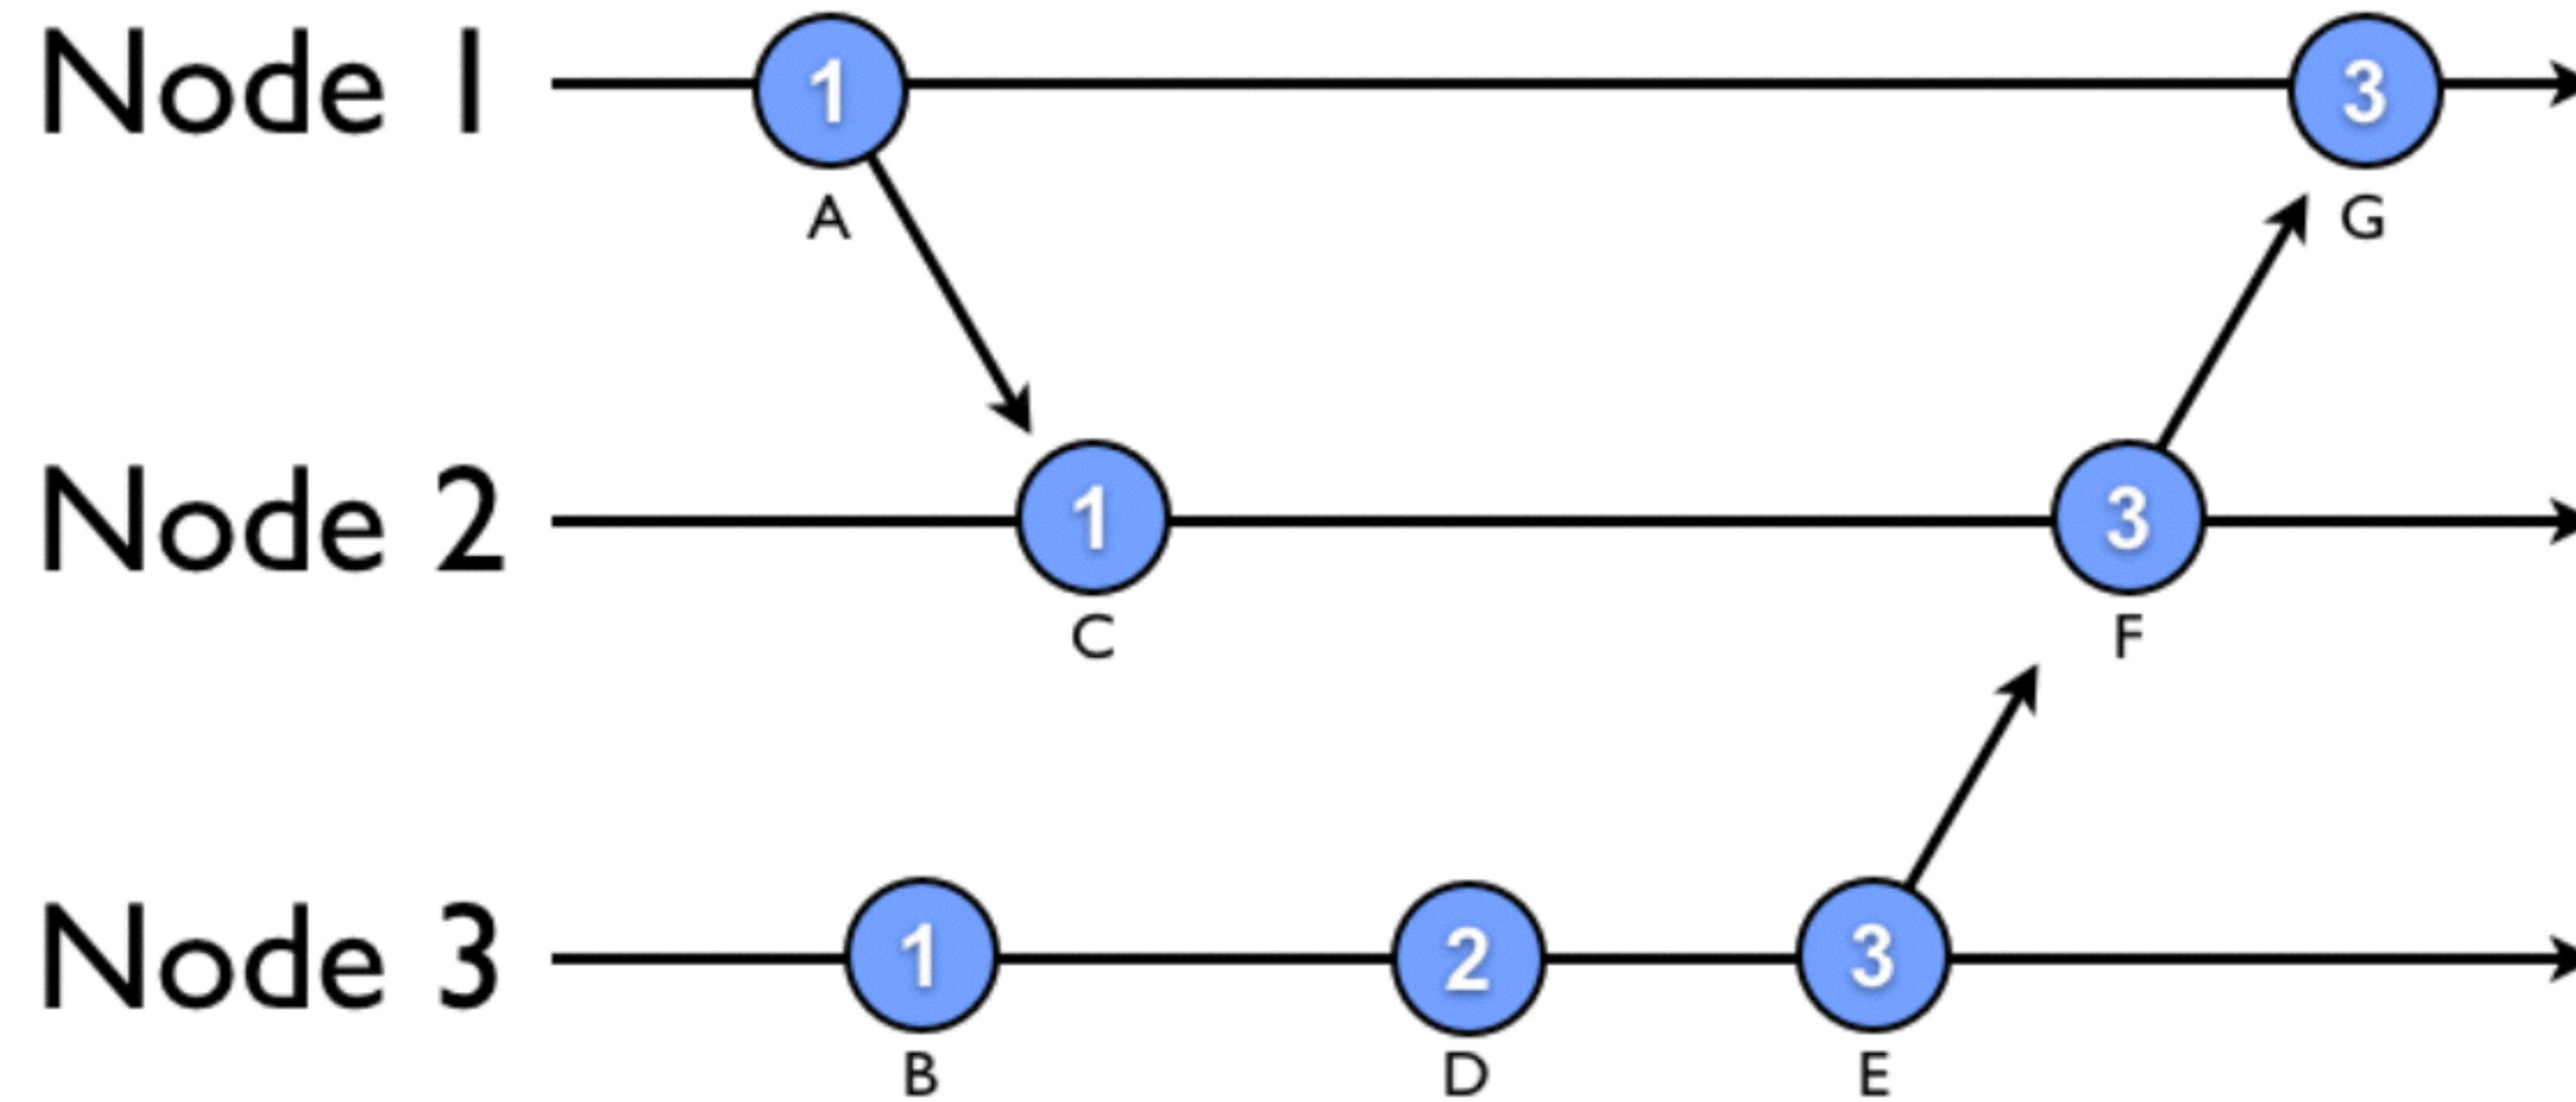
**What is a distributed system**



# Examples of nodes

- ▶ Web server
- ▶ A browser
- ▶ Instance of a mobile application
- ▶ IOT device
- ▶ A message broker
- ▶ Datastore
- ▶ Another distributed system

# Nodes in the System



# Example of communication between nodes

- ▶ Browser makes request to server (via page load/ajax)
- ▶ Server makes request to database (SQL call)
- ▶ Server makes request to another server (HTTP request)
- ▶ Datastore replication (clustering)
- ▶ Messaging (rabbitmq, kafka, etc...)



*“A distributed system is one in which the failure of a computer you didn't even know existed can render your own computer unusable”*

**- Leslie Lamport**

# Why even distribute a system?

- ▶ Scale due to performance
- ▶ Redundancy
- ▶ Scale due to data size
- ▶ Availability
- ▶ Physical distance
- ▶ Microservices\*

➡ Take advantage of the internet

## **8 Fallacies of Distributed Computing**

1-7 Peter Deutsch (1994)

8 James Gosling (1997)

# Fallacy #1

The network is reliable

- ▶ Traffic is reordered
  - ▶ Traffic is redelivered
  - ▶ Traffic is dropped or delayed
  - ▶ TCP largely handles these issues, but only per connection
- 
- ➡ Use error handling callbacks for everything network related (Promise.catch, callback err)
  - ➡ Have a plan for unavailable resources (progressive enhancement)

# Fallacy #2

Latency is zero

- ▶ Bounded by physics
  - ▶ Specifically the speed of light
  - ▶ 1 foot per nanosecond - (See Grace Hopper <https://www.youtube.com/watch?v=JEpsKnWZrJ8>)
  - ▶ Beware treating remote calls like local calls
  - ▶ Only solution is to move data closer to the consumer
- ➡ Assume any remote call can and will be slow (AJAX, HTTP, SQL, Write to Disk, etc...)

# Fallacy #3

Bandwidth is infinite

- ▶ Still an issue, but less of one than in the past
  - ▶ Still tends to show up because of mobile networks
  - ▶ Differs wildly based on geographic location
  - ▶ If you have large amounts of data this is still a very real issue
  - ▶ Amazon Snowmobile - Will truck your data from a datacenter to “the cloud”
- ➡ Don't load more data than you need (avoid unneeded JS/CSS, minify/pack your sources)

# Fallacy #4

The network is secure

- ▶ Build with security in mind, don't be complacent
  - ▶ The internet is dangerous
  - ▶ Spoofing, DDOS, MITM
  - ▶ Watch Luke Crouch's Thunderplains talk for more related things - <https://www.youtube.com/watch?v=OXDpJUhDTos>
- 
- ➡ Use HTTPS, don't send/store secure data in plain text
  - ➡ Don't trust user input

# Fallacy #5

Topology doesn't change

- ▶ Don't depend on specific routes/locations
  - ▶ DNS is a (decent) abstraction around IP locations
  - ▶ Don't assume quick paths will always stay quick
  - ▶ Nodes on the network move
- 
- ➡ Have a plan for being offline
  - ➡ Be aware of DNS, how it works, it's issues



# Fallacy #6

There is one administrator

- ▶ You don't control the network
  - ▶ VPNs, Firewalls, etc...
  - ▶ Especially prevalent for the web
  - ▶ Can you fallback to different protocols if you cannot communicate
- ➡ Provide fallbacks for things like websockets, web RTC

# Fallacy #7

Transport cost is zero

- ▶ Cost as in Money

- ▶ Highly available, low latency networks/systems can be built, but cost a lot money

- ▶ Cost as in Resources

- ▶ Transporting data over the network consumes system resources

➡ Microservices == cost

➡ Avoid additional network calls (AJAX, CSS, JS)

# Fallacy #8

The network is homogeneous

- ▶ Different parts of the network have different latency/bandwidth
  - ▶ Last leg tends to be slow
  - ▶ Some network links are less reliable than others
- 
- ➡ Calls to different servers will have different behaviors
  - ➡ Have a plan for partial availability of web resources (what happens when your web font doesn't load?)



# Plan for failure

- ▶ If possible, retry
- ▶ Utilize Idempotency
- ▶ Respect back pressure
- ▶ Circuit breakers
- ▶ Assume duplicate delivery
- ▶ Enforce or avoid explicit ordering
- ▶ Use queueing to insulate yourself from failure
- ▶ Instrument your system

Day to day problems

# Unreliable networks

- ▶ How do you handle slow networks?
- ▶ How do you communicate network problems to the user?
- ▶ How do you handle partial failure? Rollback?

# Bad user behavior

- ▶ How do you handle things like double submit problems?
- ▶ What if a user reloads a POST'd page?
- ▶ How do you handle user impatience? Can you avoid this?



# Concurrency

- ▶ How to handle multiple concurrent updates to the same resource?
- ▶ Can you enforce ordering? Should you?
- ▶ How to avoid dirty writes?

# “The app is slow”

- ▶ What is slow?
- ▶ What does slow mean?
- ▶ Where is it slow?
- ▶ Can you insulate against the slowness?

**Final thoughts**

Distributed Systems are **Challenging**

Distributed Systems are **Unavoidable**

Distributed Systems are **Necessary**

# Resources

- ▶ Contemporaries

- ▶ Caitie McCaffrey

- ▶ Camille Fournier

- ▶ Christopher Meiklejohn

- ▶ Ines Sombre

- ▶ Kyle Kingsbury

- ▶ Papers

- ▶ Distributed Systems reading list - <https://dancres.github.io/Pages/>

Questions?

---